

# COL - 774 Machine Learning (Assignment -1)

Om Prakash (2019MCS2567)

Monday, 10 Feb, 2020

## 1 Linear Regression

### 1.1 Que.1(a)

Learning Rate = 0.001

Theta 0 : 0.99661694

Theta 1 : 0.00134019

Final Cost : 1.1947947960167756e-06

No. of iterations : 12656

Stopping Criteria: 0.000000000000001: ( $10^{-14}$ ) (If the difference between the two consecutive values of error function  $J(\theta)$  is less than  $10^{-14}$ )

### 1.2 Que.1(b)

Hypothesis function is a straight line and can be written in the following form:

$y = mx + c$ , where  $m = 0.00134019$  and  $c = 0.99661694$




Figure 1: Linear Regression using  $\eta = 0.001$




Figure 2: Error Function using  $\theta_0 = 0.99661694$  and  $\theta_1 = 0.00134019$

### 1.3 Que.1(c)

3 dimensional mesh showing the error function  $J(\theta)$  on z-axis and the parameters in the x-y plane.




Figure 3: 3D Mesh grid and convergence of batch gradient descent for  $\eta = 0.25$

### 1.4 Que.1(d)

The contours of the error function at each iteration of the gradient descent. Following are the learned values:

Final Theta : [0.99661997 0.0013402 ]

Final Cost : 1.1947898199531533e-06

No. of iterations : 55




Figure 4: Contour representation of gradient descent at  $\eta = 0.25$

### 1.5 Que1.(e)

With  $\eta = 0.001$ , error function converges very slowly and decreases at a very minimal rate. Following are the learned parameters:

Final Theta : [0.99661694 0.00134019]

Final Cost : 1.1947947960167756e-06

No. of iterations : 12656




Figure 5: Contour representation of gradient descent at  $\eta = 0.001$




Figure 6: Contour representation of gradient descent at  $\eta = 0.025$

Following are the learned parameters:

Final Theta : [0.99661949 0.0013402 ]

Final Cost : 1.1947899977272803e-06

No. of iterations : 565




Figure 6: Contour representation of gradient descent at  $\eta = 0.1$

With  $\eta = 0.1$ , error function converges rapidly and decreases at a very faster rate, following are the values generated:

Final Theta : [0.99661981 0.0013402 ]

Final Cost : 1.1947898516642842e-06

No. of iterations : 143

### Conclusion:

- It has been observed if the learning rate is larger then the error function converges faster with lesser number of iterations required.
- If the learning rate is smaller then the error function converges slower and needs more number of iterations as shown above.
- If the learning rate is very high then the error function may not converge and overshoot.
- If learning parameters are too high then python may result into nan values as can't handle such big values.

## 2. Stochastic Gradient Descent

### 2.1 Que.2(a):

Sample 1 million data points:

```
x1 = np.random.normal(3,2,sample_count).reshape(100000,1)
x2 = np.random.normal(-1,2,sample_count).reshape(100000,1)
e = np.random.normal(0,math.sqrt(2),sample_count).reshape(100000,1)
```

### 2.2 Que.2(b):

Following are the learned attributes by the SGD algorithm for different batch size.

No. of samples : 1000000 Batch Size : 1 Total Iterations: 50000

Final Theta :

```
[[2.99968088]
 [0.95472704]
 [1.97118111]]
```

Execution Time in seconds : 0.88

Converge Criteria: Max\_iterations==50000 or the difference between the average cost on 1000 examples is less than  $10^{-3}$ .

No. of samples : 1000000 Batch Size : 100 Total Iterations: 20000

Final Theta :

```
[[2.98449633]
 [1.00397983]
 [1.99754537]]
```

Execution Time in seconds : 0.39

Converge Criteria: max\_iterations==100000 or the difference between the average cost on 1000 examples is less than  $10^{-3}$ .

No. of samples : 1000000 Batch Size : 10000 Total Iterations: 26000

Final Theta :

```
[[2.99101259]
 [1.0017618 ]
 [1.99969984]]
```

Execution Time in seconds : 4.09

Converge Criteria: The difference between the average cost on 1000 examples is less than  $10^{-6}$ .

No. of samples : 1000000 Batch Size : 1000000 Total Iterations: 26000

Final Theta :

[[2.99776263]

[1.00061538]

[1.9999845 ]]

Execution Time in seconds : 339.5

Converge Criteria: The difference between the average cost on 1000 examples is less than  $10^{-6}$ .

### 2.3 Que.2(c):

- a. Do different algorithms in the part above (for varying values of  $r$ ) converge to the same parameter values?

**Ans:** No, as shown above, values are different for different  $r$  values.

- b. How much different are these from the parameters of the original hypothesis from which the data was generated?

**Ans:** As shown above, the difference is very minimal, they almost converged to the original theta parameters.

- c. Comment on the relative speed of convergence and also on the number of iterations in each case?

**Ans:**

$r = 1$ , converge very fast. Max number of iterations = 50,000, Epoch = 1, though it converged even in less than 50000 iterations, it took 0.88 seconds of time for convergence.

$r = 100$ , faster than the previous one, max iterations = 20000, Epoch = 1, it took 0.39 seconds of time for convergence.

$r = 10000$ , Execution Time in seconds : 4.09, Total Iterations: 26000, No of Epochs : 260

$r = 10,00,000$ , The slowest one, Execution Time in seconds : 339.5, total iterations = 26000, No of Epochs : 26000



- d. Test error with respect to the prediction of the original hypothesis, and compare with the error obtained using learned hypothesis in each case?

**Ans:**

1. Batch Size = 1  
original cost : 0.9829469215000001  
My Theta : [2.99968088 0.95472704 1.97118111]  
My Cost : 1.1337547186188743  
Error difference : 0.15080779711887426
2. Batch Size = 100  
original cost : 0.9829469215000001  
My Theta : [2.98449633 1.00397983 1.99754537]  
My Cost : 0.9842663233575532  
Error difference : 0.001319401857553082
3. Batch Size = 10000  
original cost : 0.9829469215000001  
My Theta : [2.99101259 1.0017618 1.99969984]  
My Cost : 0.9831158377688991  
Error difference : 0.0001689162688990331
4. Batch Size = 10,00,000  
original cost : 0.9829469215000001  
My Theta : [2.99776263 1.00061538 1.9999845 ]  
My Cost : 0.9829446735402787  
Error difference : -2.2479597213687086e-06

### 3. Logistic Regression

Logistic regression transforms its output using the logistic sigmoid function to return a probability value which can then be mapped to two or more discrete classes.

Following is the log likelihood function of logistic regression:

$$L(\theta) = \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))$$

Following is the Newton's method:

$$\theta^{(t+1)} = \theta^{(t)} - H^{-1} \nabla_{\theta} L(\theta)$$

#### 3.1 Que.3(a):

Following are the values of theta which are learned by the algo:

```
[[ 0.40125316]
 [ 2.5885477]
 [-2.72558849]]
```

#### 3.2 Que.3(b):

A straight line showing the boundary separating the region where  $h(x) > 0.5$  from where  $h(x) \leq 0.5$ .




Figure 7: Logistic Regression using Newton's method

#### 4. Gaussian Discriminant Analysis (GDA)

The model is:

$$\begin{aligned}y &\sim \text{Bernoulli}(\phi) \\x|y=0 &\sim \mathcal{N}(\mu_0, \Sigma) \\x|y=1 &\sim \mathcal{N}(\mu_1, \Sigma)\end{aligned}$$

Writing out the distributions, this is:

$$\begin{aligned}p(y) &= \phi^y(1-\phi)^{1-y} \\p(x|y=0) &= \frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x-\mu_0)^T\Sigma^{-1}(x-\mu_0)\right) \\p(x|y=1) &= \frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x-\mu_1)^T\Sigma^{-1}(x-\mu_1)\right)\end{aligned}$$

The log-likelihood of the data is given by

$$\begin{aligned}\ell(\phi, \mu_0, \mu_1, \Sigma) &= \log \prod_{i=1}^m p(x^{(i)}, y^{(i)}; \phi, \mu_0, \mu_1, \Sigma) \\&= \log \prod_{i=1}^m p(x^{(i)}|y^{(i)}; \mu_0, \mu_1, \Sigma) p(y^{(i)}; \phi).\end{aligned}$$

By maximizing  $\ell$  with respect to the parameters, we find the maximum likelihood estimate of the parameters to be:

$$\begin{aligned}\phi &= \frac{1}{m} \sum_{i=1}^m 1\{y^{(i)} = 1\} \\\mu_0 &= \frac{\sum_{i=1}^m 1\{y^{(i)} = 0\} x^{(i)}}{\sum_{i=1}^m 1\{y^{(i)} = 0\}} \\\mu_1 &= \frac{\sum_{i=1}^m 1\{y^{(i)} = 1\} x^{(i)}}{\sum_{i=1}^m 1\{y^{(i)} = 1\}} \\\Sigma &= \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{y^{(i)}})(x^{(i)} - \mu_{y^{(i)}})^T.\end{aligned}$$

Note that covariance matrix  $\Sigma$  is identical here.

#### 4.1 Que.4(a)

Both the classes have the same covariance matrix i.e.  $\Sigma_0 = \Sigma_1 = \Sigma$ .

$$\begin{aligned}\phi &= \frac{1}{m} \sum_{i=1}^m 1\{y^{(i)} = 1\} \\ \mu_0 &= \frac{\sum_{i=1}^m 1\{y^{(i)} = 0\} x^{(i)}}{\sum_{i=1}^m 1\{y^{(i)} = 0\}} \\ \mu_1 &= \frac{\sum_{i=1}^m 1\{y^{(i)} = 1\} x^{(i)}}{\sum_{i=1}^m 1\{y^{(i)} = 1\}} \\ \Sigma &= \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu_{y^{(i)}})(x^{(i)} - \mu_{y^{(i)}})^T.\end{aligned}$$

$\mu_0$  : [-0.75529433 0.68509431]

$\mu_1$  : [ 0.75529433 -0.68509431]

Sigma ( $\Sigma$ ) :

[[ 0.42953048 -0.02247228]

[-0.02247228 0.53064579]]

#### 4.2 Que.4(b)

The training data corresponding to the two coordinates of the input features, Note that Alaska is considered 0 here and Canada is considered as 1.

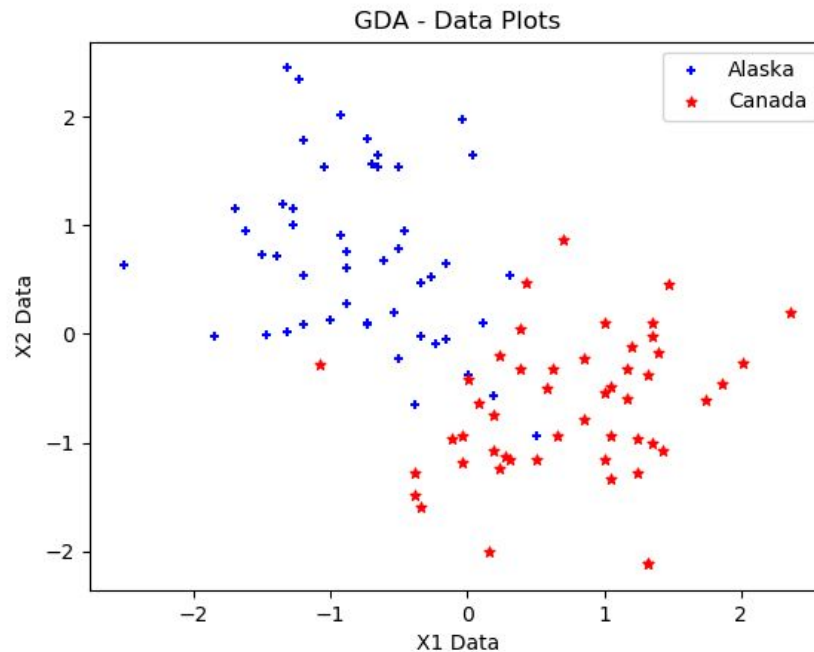


Figure 8: Training data corresponding to the two coordinates of the input features

#### 4.3 Que.4(c)

Equation of the boundary separating the two regions in terms of the parameters  $\mu_0$ ,  $\mu_1$  and  $\Sigma$ , when the two classes have identical covariance matrix. Following is the equation of boundary when the two covariance matrices are the same.

$$(\mu_0^T - \mu_1^T) \Sigma^{-1} X - \left( \log\left(\frac{\phi}{1-\phi}\right) + \frac{1}{2}(\mu_1^T \Sigma^{-1} \mu_1 - \mu_0^T \Sigma^{-1} \mu_0) \right) = 0$$

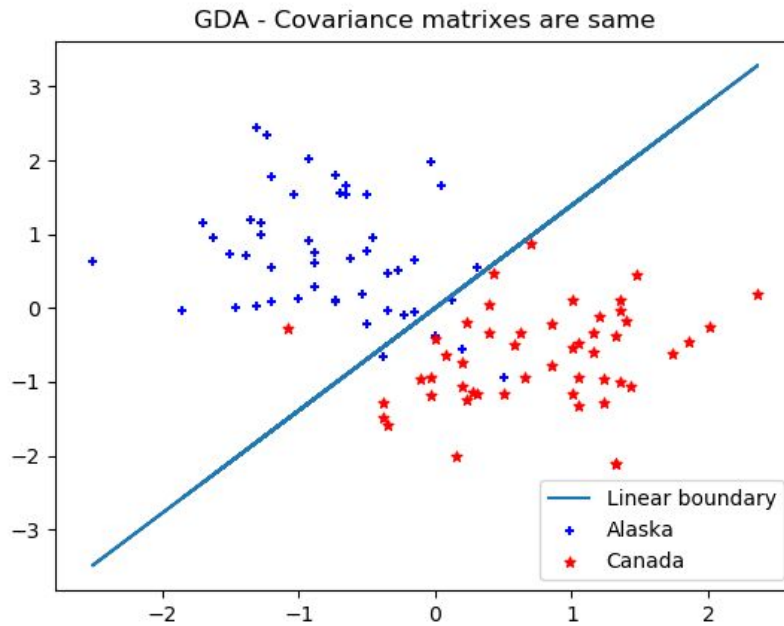


Figure 9: Linear separator when the two classes have identical covariance matrix

#### 4.4 Que.4(d)

Following are the different parameters learned by the algorithm when the two covariance  $\Sigma_0$  and  $\Sigma_1$  matrices are the same.

$$\Sigma_0 : \begin{bmatrix} 0.38158978 & -0.15486516 \\ -0.15486516 & 0.64773717 \end{bmatrix}$$

$$\Sigma_1 : \begin{bmatrix} 0.47747117 & 0.1099206 \\ 0.1099206 & 0.41355441 \end{bmatrix}$$

$$\mu_0 : [-0.75529433 \quad 0.68509431]$$

$$\mu_1 : [0.75529433 \quad -0.68509431]$$

#### 4.5 Que.4(e)

The maximum-likelihood estimate of the covariance matrix  $\Sigma_0$  can be derived using the equation:

$$\Sigma_0 = \frac{\sum_{i=1}^m 1\{y^{(i)} = 0\} (x^{(i)} - \mu_{y^{(i)}})(x^{(i)} - \mu_{y^{(i)}})^T}{\sum_{i=1}^m 1\{y^{(i)} = 0\}}$$

It's the same for  $\Sigma_1$  but  $\Sigma_0$  will be replaced with  $\Sigma_1$ .

The equation for the quadratic boundary separating the two regions:  $\Sigma_1$

$$\frac{1}{2} X^T (\Sigma_1^{-1} - \Sigma_0^{-1}) X + (\mu_0^T \Sigma_0^{-1} - \mu_1^T \Sigma_1^{-1}) X - \left( \log\left(\frac{\phi}{1-\phi}\right) + \frac{1}{2} \log\left(\frac{|\Sigma_0|}{|\Sigma_1|}\right) \right) + \frac{1}{2} (\mu_1^T \Sigma_1^{-1} \mu_1 - \mu_0^T \Sigma_0^{-1} \mu_0) = 0$$





Figure 10: Quadratic boundary the covariance matrix is different

#### 4.6 Que.4(f)



Note that the two Gaussians have contours that are the same shape and orientation, since they share a covariance matrix  $\Sigma$ , but they have different means  $\mu_0$  and  $\mu_1$ . Also shown in the figure is the straight line giving the decision boundary at which  $p(y = 1|x) = 0.5$ . On one side of the boundary, we'll predict  $y = 1$  to be the most likely outcome, and on the other side, we'll predict  $y = 0$ .

The quadratic boundary by GDA offers a non linear way of classifying data. Although it computes an additional sigma parameter, it provides a better estimate of the decision boundaries.

**Reference:** Andrew NG Notes